



Data Analysis & Visualisation

CSC3062

BEng (CS & SE), MEng (CS & SE), BIT & CIT

Dr Reza Rafiee

Semester 1 2019



Supervised learning



What we need to know about classification

- What is classification?
- What we need as a dataset in classification
- Binary vs. multiclass classification
- Classification models (categories of classifier models)
- How to choose a classification model?
- Support vector machine (SVM) classifier model
- Designing a multiclass SVM model with an example
- How to evaluate the performance of a classifier model?


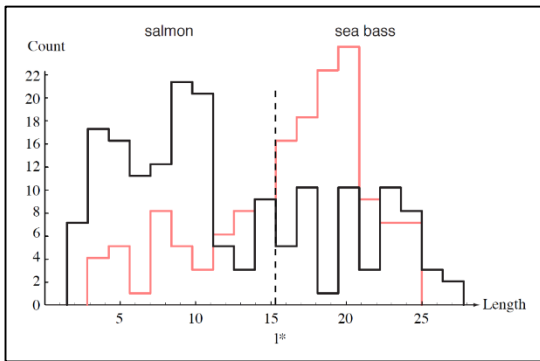
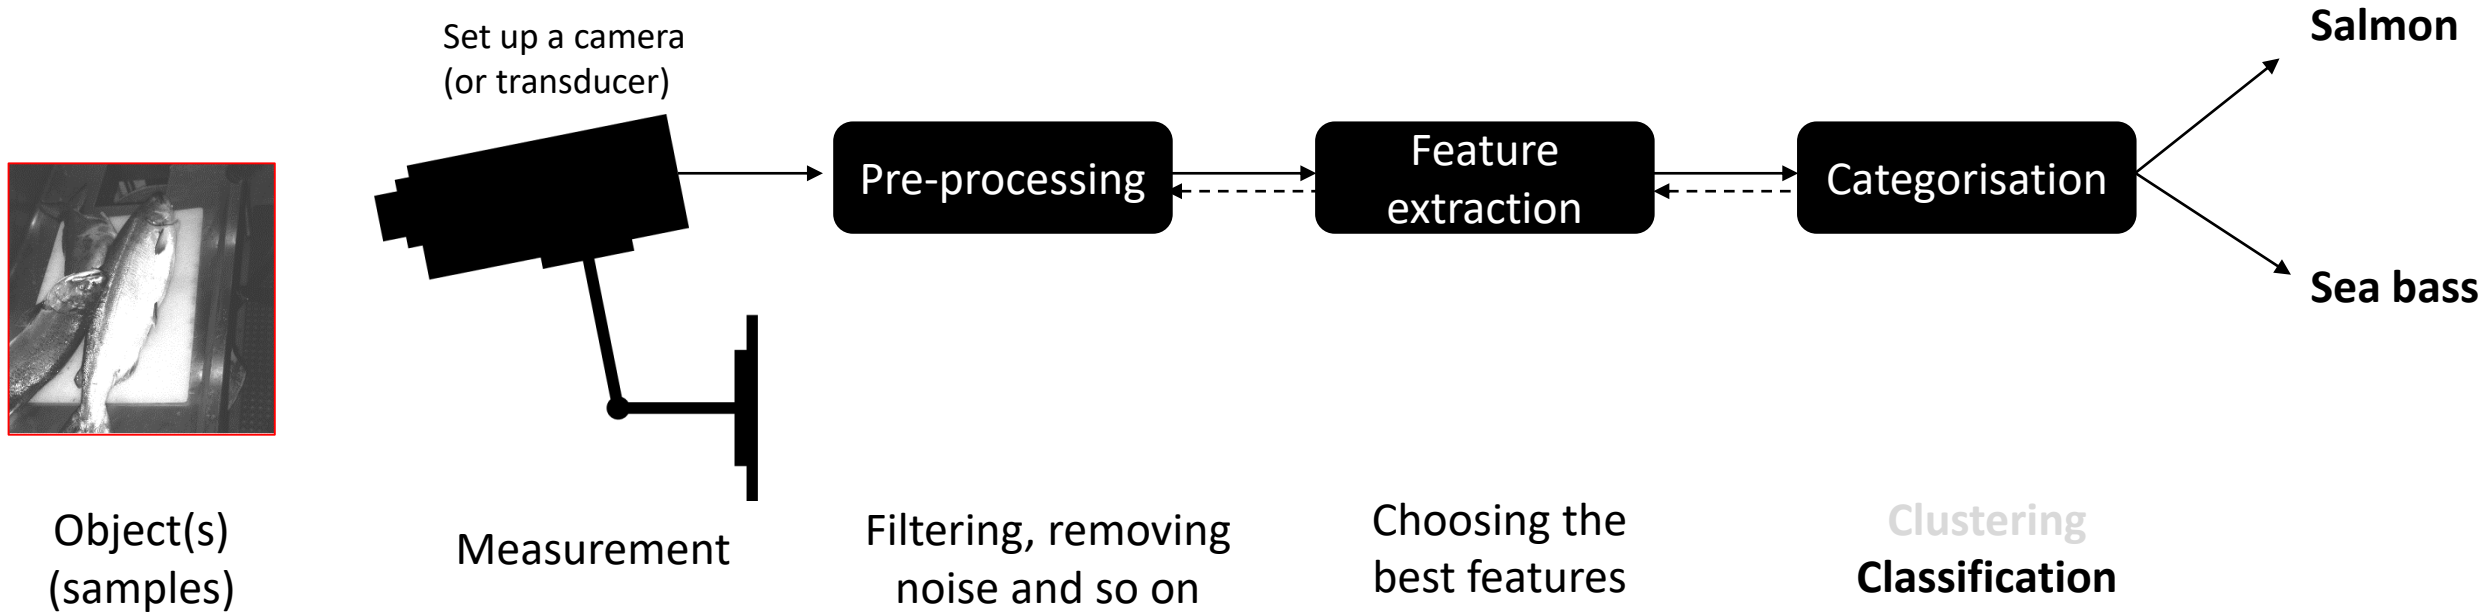


Classification applications

- Text categorisation (e.g., spam filtering)
- Optical character recognition (OCR) (e.g., a computerised system transferring hard documents into word doc.)
- Machine vision (e.g., face detection)
- Natural-language processing (NLP) (e.g., spoken language understanding)
- Market segmentation (e.g., predict if customer will respond to promotion)
- Bioinformatics and medicine (e.g., classify cancer patients into different immune subtypes)
- ...



Pattern recognition systems



x_1 : lightness
 x_2 : width


The aim is to partition the feature space into two regions

Feature space: two dimensions $X = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$

Suppose that we measure the feature vectors for our samples



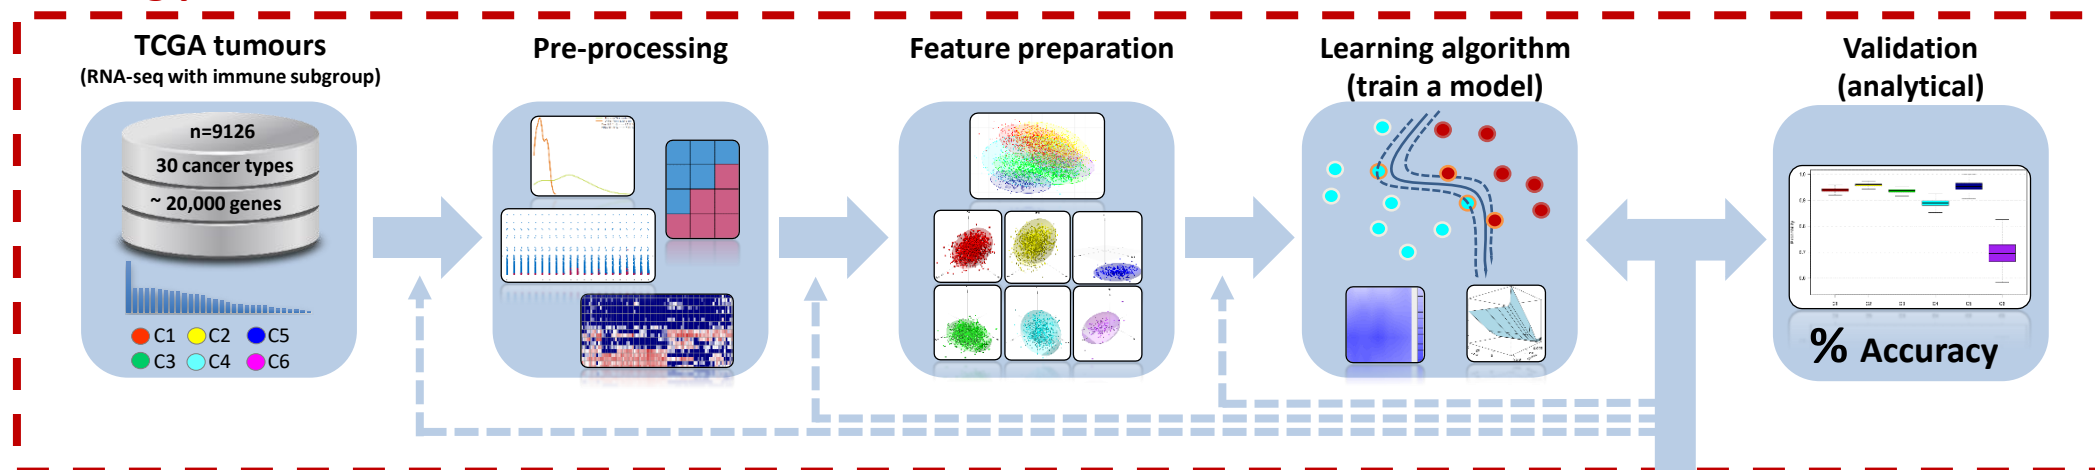
Classification vs. clustering



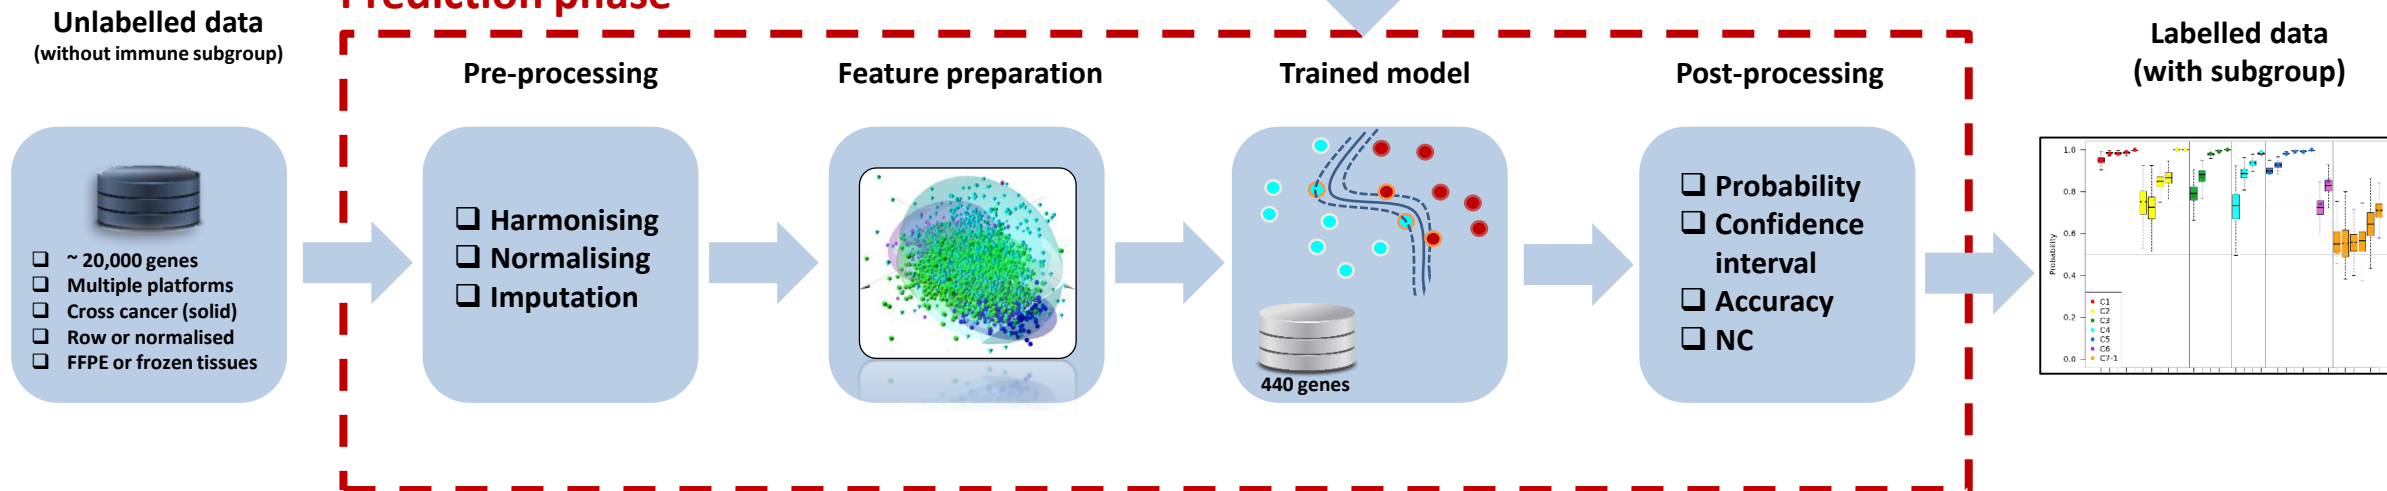


Classification – training vs. prediction

Training phase



Prediction phase



Classification stages. a, Training phase. b, Prediction phase.



Prediction problem

Consider the following example

It is important for the bank to be able to **predict** in advance **the risk associated with a loan**, which is the probability that the customer will default and not pay the whole amount back.

In **credit scoring**, the bank calculates the risk given the amount of credit and **the information about the customer**. The **information about the customer** includes data we have access to and is relevant in calculating his or her financial capacity - namely, **income, savings, collaterals, profession, age, past financial history**, and so forth. The bank has a record of past loans containing such customer data and whether the loan was paid back or not.

From this data of particular applications, the aim is **to infer a general rule** coding **the association between a customer's attributes and his risk**.

A machine learning system fits a model to the past data to be able to calculate the risk for a new application and then decides to accept or refuse it accordingly.




Prediction problem - classification

- It is important for the bank to be able to **predict** in advance **the risk associated with a loan**, which is the probability that the customer will default and not pay the whole amount back.
- From a data of particular applications, the aim is to infer a general rule coding **association** between a **customer's attribute (features)** and **his/her risk**.
- This is an example of a ***classification*** problem where there are two classes: low-risk and high-risk customers. The information about a customer makes up the ***input*** to the classifier whose task is to assign the input to one of the two classes.




Dataset in classification


Features



17 samples
(customers)




#	Income	Savings	Collaterals	Profession	age	Risk
1	£25K	2K	Yes	...	24	low
2	£35K	5K	No	...	30	low
3	£5K	0.5K	Yes	...	26	high
...
17	£45K	10K	No	...	40	low



Class labels
or
subgroups



Dataset in classification



This figure illustrates an example of a dataset. Each circle corresponds to one **data instance** with input values in the corresponding axes. **For simplicity**, only two customer **attributes or features**, income and savings, are taken **as input** and the two classes are low-risk ('+') and high-risk ('-').



Prediction problem – *discriminant*

X_1 : Income

X_2 : Savings

Y : low-risk or high-risk




- This is an example of *discriminant*, it is a function that separates the examples of different classes.
- **Discriminant analysis**



Prediction problem; parameters of the model

Model
parameters:
 θ_1 & θ_2



IF $x_1 > \theta_1$ and $x_2 > \theta_2$ THEN low-risk ELSE high-risk

Once we have a rule like this that fits the past data, if the future is similar to past, then we can make correct predications for new instances.

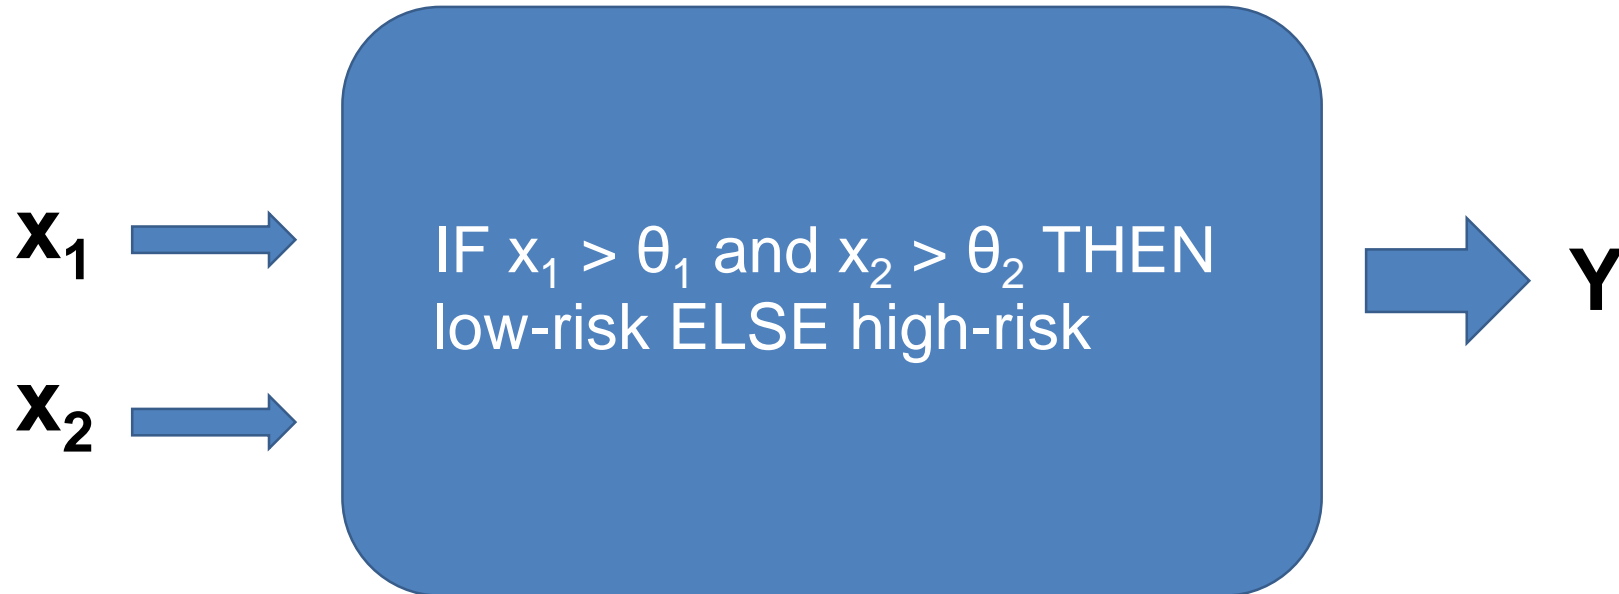


Prediction problem – *discriminant*

X_1 : Income

X_2 : Savings

Y (output): low-risk or high-risk



- This is an example of *discriminant*, it is a function that separates the examples of different classes.
- Discriminant analysis



Classification - concept

Basic idea: classify samples into
pre-defined groups

What is difference between classification and regression?



Regression vs. classification

Basic idea: classify samples into
pre-defined groups

A regression algorithm predicts a **continuous** output

A classification algorithm predicts a **discrete** output (class label)



Prediction in regression & classification



What is the temperature going to be tomorrow?

Based on previous data

35°




Will it be cold or hot weather?

Based on previous data

Cold



Hot





Binary vs. multiclass classification

Binary classifier classifies data points into one of two classes



Multiclass classifier: classifies data points into one of three or more classes





Classification algorithms

- ☐ K-Nearest Neighbour
- ☐ Naive Bayes Classifier
- ☐ Support Vector Machines (the basic SVM supports only binary classification); linear or with Gaussian kernels
- ☐ Decision Trees (e.g., Random Forest)
- ☐ Artificial Neural Networks (ANN)
- ☐ Hierarchical classifier
- ☐ ...



Classification algorithms

- ☐ K-Nearest Neighbour
- ☐ Naive Bayes Classifier
- ☐ **Support Vector Machines** (the basic SVM supports only binary classification); linear or **with Gaussian kernels**
- ☐ Decision Trees (e.g., Random Forest)
- ☐ Artificial Neural Networks (ANN)
- ☐ Hierarchical classifier
- ☐ ...



Parametric vs. nonparametric models

In **parametric** “learners” or models, we have a fixed size of parameters (the weight coefficient). This makes them not very flexible.

- ☐ Linear Regression
- ☐ Linear Support Vector Machines
- ☐ Logistic Regression
- ☐ Naive Bayes



Parametric vs. nonparametric models

While **nonparametric** “learners” or models are good when you have a large data and you don’t want to worry too much about choosing just the right features

- ☐ Decision Trees
- ☐ K-Nearest Neighbour
- ☐ Support Vector Machines with Gaussian Kernels
- ☐ Artificial Neural Networks

Non-parametric models assume that the data distribution cannot be defined in terms of such a finite set of parameters. This makes them more flexible



Parametric vs. nonparametric models

Linear regression

Naive Bayes

Linear SVMs

Logistic
regression

Less flexibility

Decision Trees

KNN

SVMs (nonlinear
kernels)





ANNs

More flexibility



Good vs. bad classifiers?


Assume, we have four classifiers developed on four training datasets.






Good vs. bad classifiers

Good




Sufficient data
Low training error
Simple classifier


Bad



Insufficient data



Training error too
high



Classifier too
complex



Any Questions?