



Data Analysis & Visualisation

CSC3062

BEng (CS & SE), MEng (CS & SE), BIT & CIT

Dr Reza Rafiee

Semester 1 2019



QUEEN'S
UNIVERSITY
BELFAST

SCHOOL OF
ELECTRONICS,
ELECTRICAL
ENGINEERING AND
COMPUTER SCIENCE

Supervised learning | classification

Supervised learning

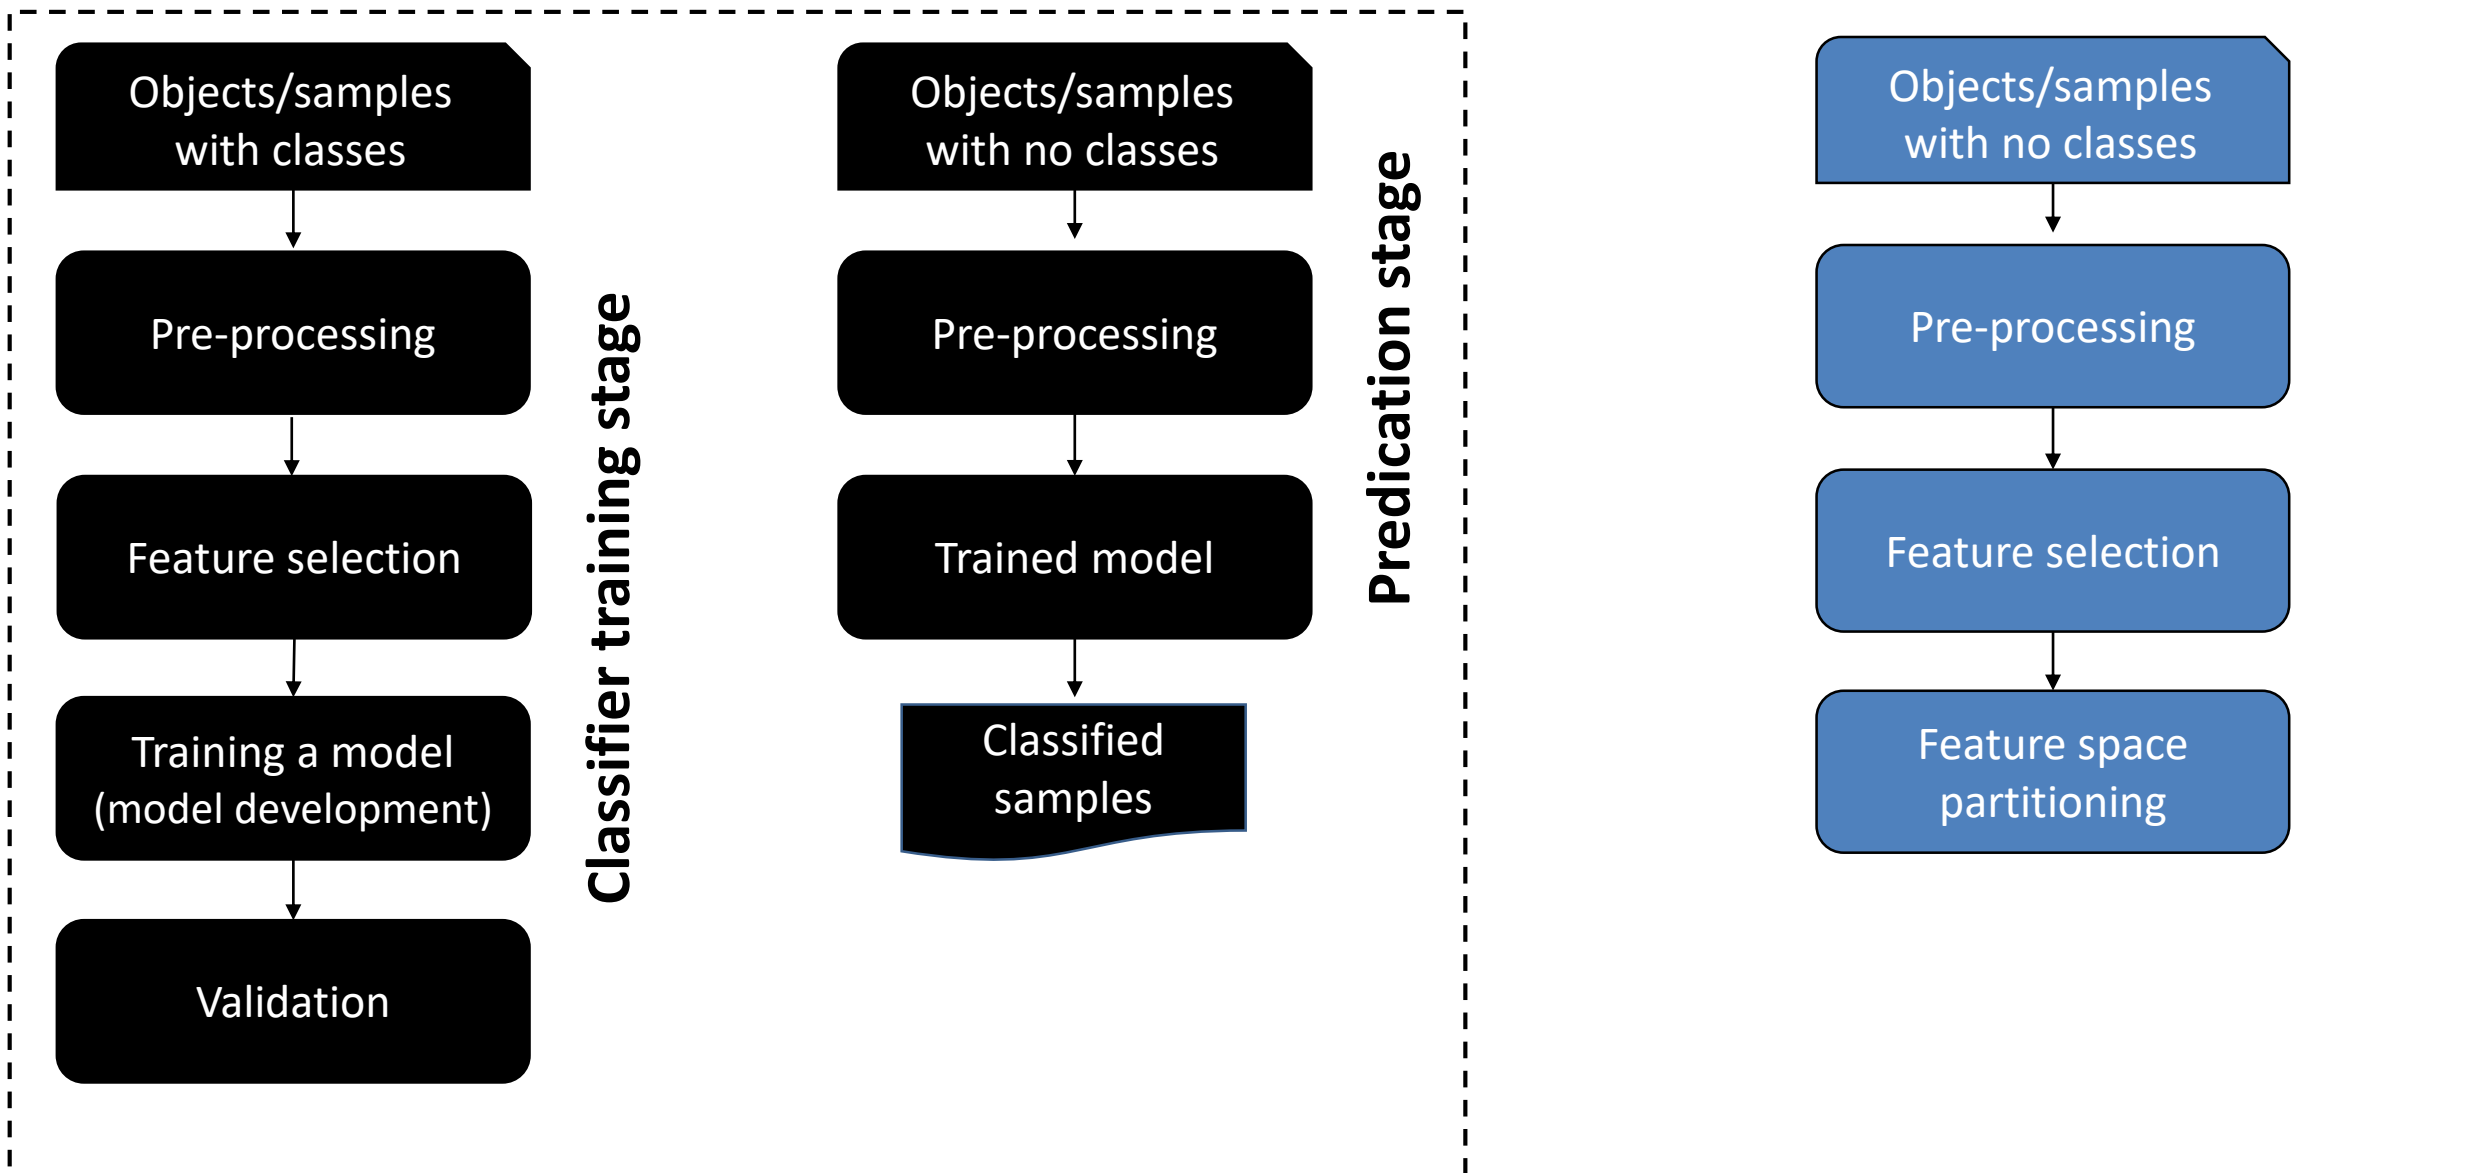


What we need to know about classification

- What is classification?
- What we need as a dataset in classification
- Binary vs. multiclass classification
- Classification models (categories of classifier models)
- How to choose a classification model?
- Support vector machine (SVM) classifier model
- Designing a multiclass SVM model with an example
- How to evaluate the performance of a classifier model?



Classification vs. clustering



Feature selection for high separability



Classification algorithms

- ☐ K-Nearest Neighbour
- ☐ Naive Bayes Classifier
- ☐ **Support Vector Machines** (the basic SVM supports only binary classification); linear or **with Gaussian kernels**
- ☐ Decision Trees (e.g., Random Forest)
- ☐ Artificial Neural Networks (ANN)
- ☐ Hierarchical classifier
- ☐ ...



Parametric vs. nonparametric models

Linear regression

Naive Bayes

Linear SVMs

Logistic
regression

Less flexibility

Decision Trees

KNN

SVMs (nonlinear
kernels)


ANNs

More flexibility




Good vs. bad classifiers

Good




Sufficient data
Low training error
Simple classifier


Bad



Insufficient data




Training error too
high



Classifier too
complex



Support vector machines (SVM)






A linear separator

Given a set of training samples, an SVM training algorithm builds a model that assigns new samples to one of the two classes (binary classifier).

Training samples
Two classes: +1 & -1




$$f(x) = \text{sign}(w^T x + b)$$

$$W^T X + b > 0 \quad \text{for red samples}$$
$$W^T X + b < 0 \quad \text{for blue samples}$$



Linear separators

Which of the linear separators is optimal?






Good separation using support vectors

Binary classification can be viewed as the task of separating classes in feature space.

A good separation is attained by the **hyperplane** that has the largest distance to the **nearest training data points of any class** (so-called functional margin), since in general the **larger the margin**, the **lower the generalisation error** of the classifier



The support vectors are indicated by the circles around them.  




Maximum margin classification

Margin λ of the separator is the distance between support vectors.

This **maximum-margin separator** is determined by a subset of the data points in a training set (“support vectors”).

In SVM, we aim to find a **right hyperplane** and then **maximize the margin** (λ) to obtain the parameters of the hyperplane (i.e., optimization problem)




The support vectors are indicated by the circles around them.  



Two key points when designing an SVM

- 1) Assess the level of your dataset complexity.
Do you need a **linear** or **non-linear/Kernel** hyperplane function as a separator?
- 2) Find a right hyperplane and then maximize the margin (λ) to obtain **the parameters of the hyperplane** (i.e., optimization problem)



The support vectors are indicated by the circles around them.  



Performance of SVM in general

- SVMs work very well in practice.
 - You must choose a linear or kernel function (i.e., hyperplane) and its parameters, but the rest is automatic.
 - The test performance is **very good**.
- SVM can be **computationally expensive** for big datasets
 - The computation of the maximum-margin hyperplane depends on the **square** of the number of training samples.



An optimal SVM classifier in R using e1071 package

1) TUNING:

```
Tuning_model <- tune(svm, Trainingset450k17, label_vector,  
scale = F, tolerance = 0.00001, type = "C-classification",  
kernel = "radial", probability = T  
ranges = list(cost= seq(0.0, 1.0, 0.2), gamma = seq(0, 15, 1)),  
tunecontrol= tune.control(sampling = "cross", cross=10), seed=123456)
```

The darkest shades of blue indicating the best (see the two plots).

Narrowing in on the darkest blue range and performing further tuning.

```
Plot(Tuning_model, xlim=range(0:15), ylim=range(0:1))
```

```
Plot(Tuning_model, xlim=range(0.2:0.25), ylim=range(8:12))
```

2) TRAINING:

```
Radial_model <- svm(Trainingset450k17, label_vector, scale  
= F, tolerance = 0.00001, type = "C-classification",  
kernel = "radial",  
cost = optimum_cost, gamma = optimum_gamma,  
probability = T, seed = 123456)
```

3) TESTING (PREDICTION):

```
Radial_model <- predict(object= Radial_model, newdata = seq_test_BEM_97, probability=T)
```

Three key steps

1) Tuning

Choose a hyperplane; try linear or nonlinear (polynomial or RBF kernels) and find it's parameters

2) Training

Train the classifier based on the identified parameters of the hyperplane

3) Testing

Test the trained classifier by giving it some new samples (without subgroups)



Find the parameters of a non-linear function (kernel function)

TUNING:

```
Tuning_model <- tune(svm, Trainingset450k17, label_vector,  
scale = F, tolerance = 0.00001, type = "C-classification",  
kernel = "radial", probability = T  
ranges = list(cost= seq(0.0, 1.0, 0.2), gamma = seq(0, 15, 1)),  
tunecontrol= tune.control(sampling = "cross", cross=10), seed=123456)
```

Input training dataset: **Trainingset450k17**

Label_vector: a vector of all sample **class labels (subgroup labels)**

1) Tuning

Choose a hyperplane and find it's parameters: **radial basis function** with two parameters which are **cost** and **gamma**

Using a **grid search**
and **10-fold cross validation technique**

Run multiple times the **tune()** to find the best (optimum) parameters




Tuning the model; grid search and 10-fold cross validation

TUNING:

```
Tuning_model <- tune(svm, Trainingset450k17, label_vector,  
scale = F, tolerance = 0.00001, type = "C-classification",  
kernel = "radial", probability = T  
ranges = list(cost= seq(0.0, 1.0, 0.2), gamma = seq(0, 15, 1)),  
tunecontrol= tune.control(sampling = "cross", cross=10), seed=123456)  
  
Plot(Tuning_model, xlim=range(0:15), ylim=range(0:1))
```

The darkest shades of blue indicating the best (see the plot).

Performance of SVM model – error rate



0 1
interval
↓
cost= seq(0.0, 1.0, 0.2)


0 15
interval
↓
gamma = seq(0, 15, 1)



Further tuning

The darkest shades of blue indicating the lowest error.

Performance of SVM model – error rate



Narrowing in on the darkest blue range and performing further tuning.



Tuning the model; grid search and 10-fold cross validation

TUNING:

```
Tuning_model <- tune(svm, Trainingset450k17, label_vector,  
scale = F, tolerance = 0.00001, type = "C-classification",  
kernel = "radial", probability = T  
ranges = list(cost= seq(8, 12, 1), gamma = seq(0.20, 0.25, 0.01)),  
tunecontrol= tune.control(sampling = "cross", cross=10), seed=123456)
```


```
Plot(Tuning_model, xlim=range(0:15), ylim=range(0:1))
```

```
Plot(Tuning_model, xlim=range(0.2:0.25), ylim=range(8:12))
```

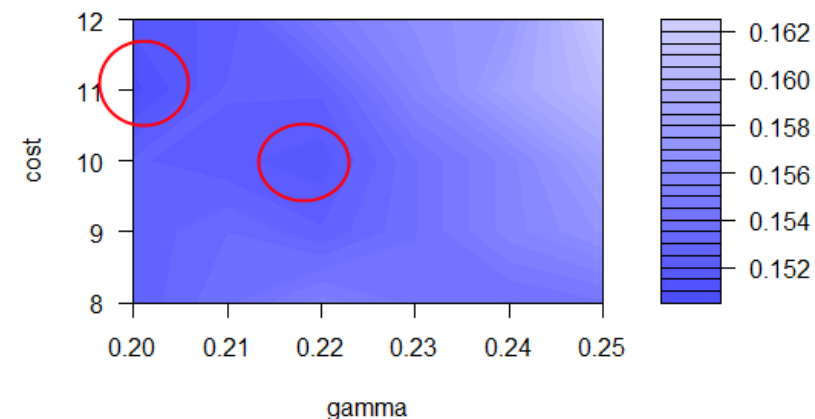
The darkest shades of blue indicating the best (see the two plots).

Narrowing in on the darkest blue range and performing further tuning.

Performance of SVM model – error rate



Performance of SVM model – error rate



8 12
┌──────────────────┐
└──────────────────┘
cost= seq(8, 12, 1)

interval
↓


0.20 0.25
┌──────────────────────────────────────────┐
└──────────────────────────────────────────┘

interval
↓
gamma = seq(0.20, 0.25, 0.01)



Final parameters of the kernel function

Performance of SVM model – error rate



gamma = 0.22
cost = 10

8 12
└────────────────┘
interval
↓
cost = seq(8, 12, 1)

0.20 0.25
└──────────────────────────────────┘
interval
↓
gamma = seq(0.20, 0.25, 0.01)



Tuning the model; grid search and 10-fold cross validation

TUNING:

```
Tuning_model <- tune(svm, Trainingset450k17, label_vector,  
scale = F, tolerance = 0.00001, type = "C-classification",  
kernel = "radial", probability = T  
ranges = list(cost= seq(8, 12, 1), gamma = seq(0.20, 0.25, 0.01)),  
tunecontrol= tune.control(sampling = "cross", cross=10), seed=123456)
```

```
Plot(Tuning_model, xlim=range(0:15), ylim=range(0:1))
```

```
Plot(Tuning_model, xlim=range(0.2:0.25), ylim=range(8:12))
```

The darkest shades of blue indicating the best (see the two plots).
Narrowing in on the darkest blue range and performing further tuning.

2) TRAINING:

```
Radial_model <- svm(Trainingset450k17, label_vector, scale = F,  
tolerance = 0.00001, type = "C-classification",  
kernel = "radial",  
cost = 10, gamma = 0.22,  
probability = T, seed = 123456)
```

Three key steps

1) Tuning

Choose a hyperplane; try linear or nonlinear (polynomial or RBF kernels) and find it's parameters

2) Training

Train the classifier based on the identified parameters of the hyperplane

3) Testing

Test the trained classifier by giving it some new samples (without subgroups)



Tuning the model; grid search and 10-fold cross validation

TUNING:

```
Tuning_model <- tune(svm, Trainingset450k17, label_vector,  
scale = F, tolerance = 0.00001, type = "C-classification",  
kernel = "radial", probability = T  
ranges = list(cost= seq(8, 12, 1), gamma = seq(0.20, 0.25, 0.01)),  
tunecontrol= tune.control(sampling = "cross", cross=10), seed=123456)
```

```
Plot(Tuning_model, xlim=range(0:15), ylim=range(0:1))
```

```
Plot(Tuning_model, xlim=range(0.2:0.25), ylim=range(8:12))
```

The darkest shades of blue indicating the best (see the two plots).

Narrowing in on the darkest blue range and performing further tuning.

2) TRAINING:

```
Radial_model <- svm(Trainingset450k17, label_vector, scale = F,  
tolerance = 0.00001, type = "C-classification",  
kernel = "radial",  
cost = 10, gamma = 0.22,  
probability = T, seed = 123456)
```

3) TESTING (PREDICTION):

```
Radial_model <- predict(object= Radial_model, newdata = seq_test_BEM_97, probability=T)
```

Three key steps

1) Tuning

Choose a hyperplane; try linear or nonlinear (polynomial or RBF kernels) and find it's parameters

2) Training

Train the classifier based on the identified parameters of the hyperplane



3) Testing

Test the trained classifier by giving it some new samples (without subgroups):
seq_test_BEM_97




Degree of polynomial features


Two groups of data



X^1



X^2



X^3

Degree of polynomial features when designing a kernel for an SVM classifier

What is resampling technique?

If you use the entire training data to select the “optimal” classifier, then there would be a fundamental problem.

The final model will normally **overfit** the training data: it will not be able to generalise to new data.

The error rate estimate will be overly optimistic (lower than the true error rate)

Split dataset into two groups

Training set: used to train the classifier

Test set: used to estimate the error rate of the trained classifier





K-fold cross-validation (CV)

Cross validation and bootstrapping are resampling methods

Question: why do we need resampling method?

**A limited number of good samples
(limited data)**

Collection of data is expensive

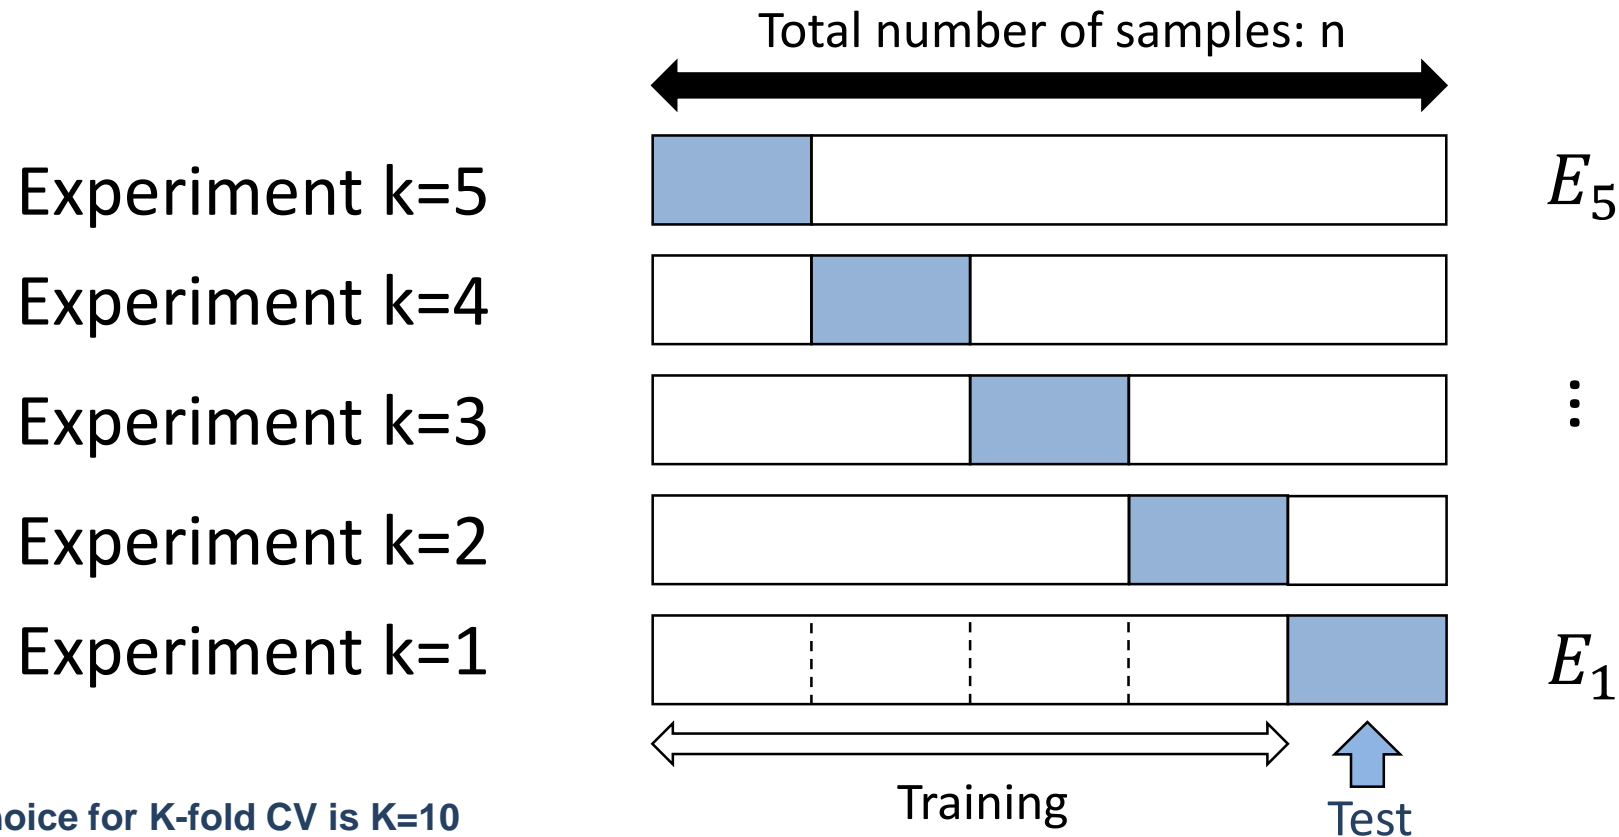


K-fold cross-validation (CV)

Create a K-fold partition of a dataset

For each of K experiments, use K-1 folds for training and a different fold for testing

This procedure is illustrated in the following figure for K=5





$$E = \frac{1}{K} \sum_{i=1}^K E_i$$

Average error



Binary vs. multiclass classification

Binary classifier classifies data points into one of two classes



Multiclass classifier: classifies data points into one of three or more classes








Multiclass to binary classification

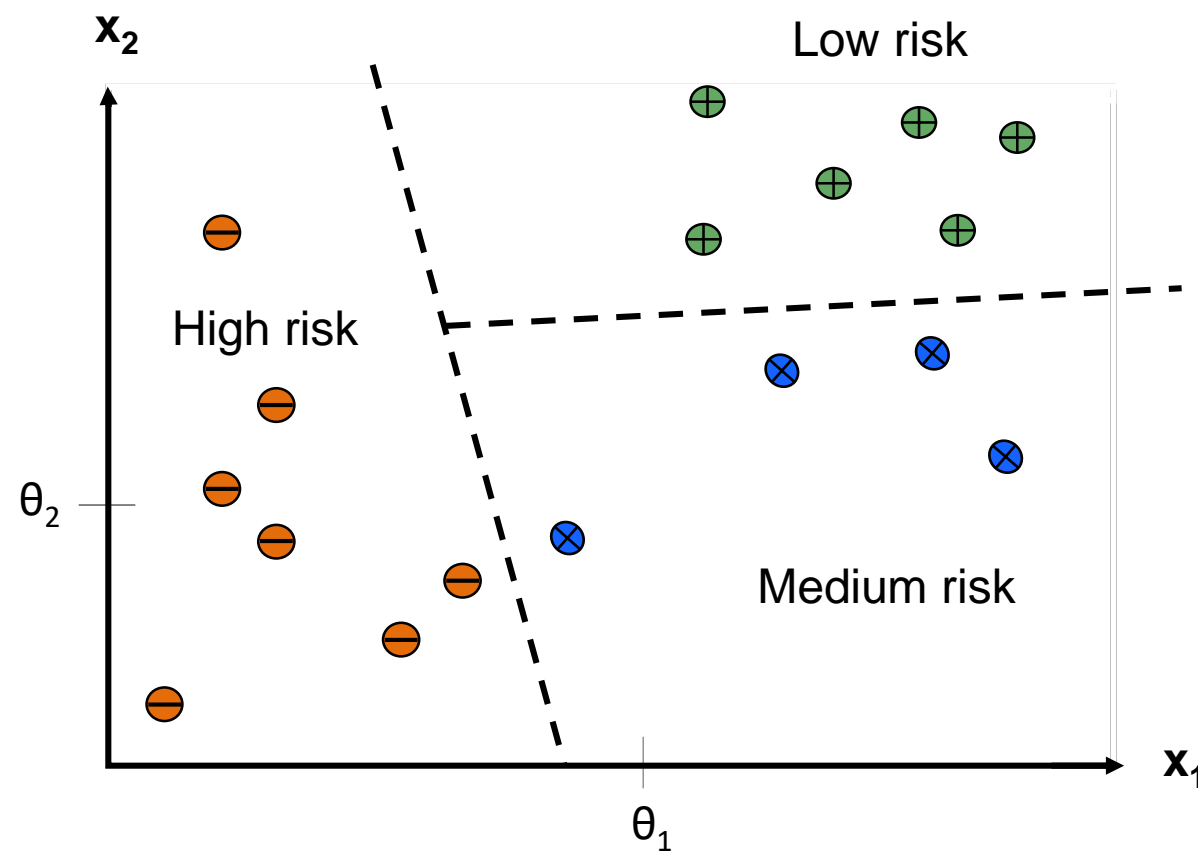
High risk: H

Medium risk: M

Low risk: L



One vs. rest (all) approach

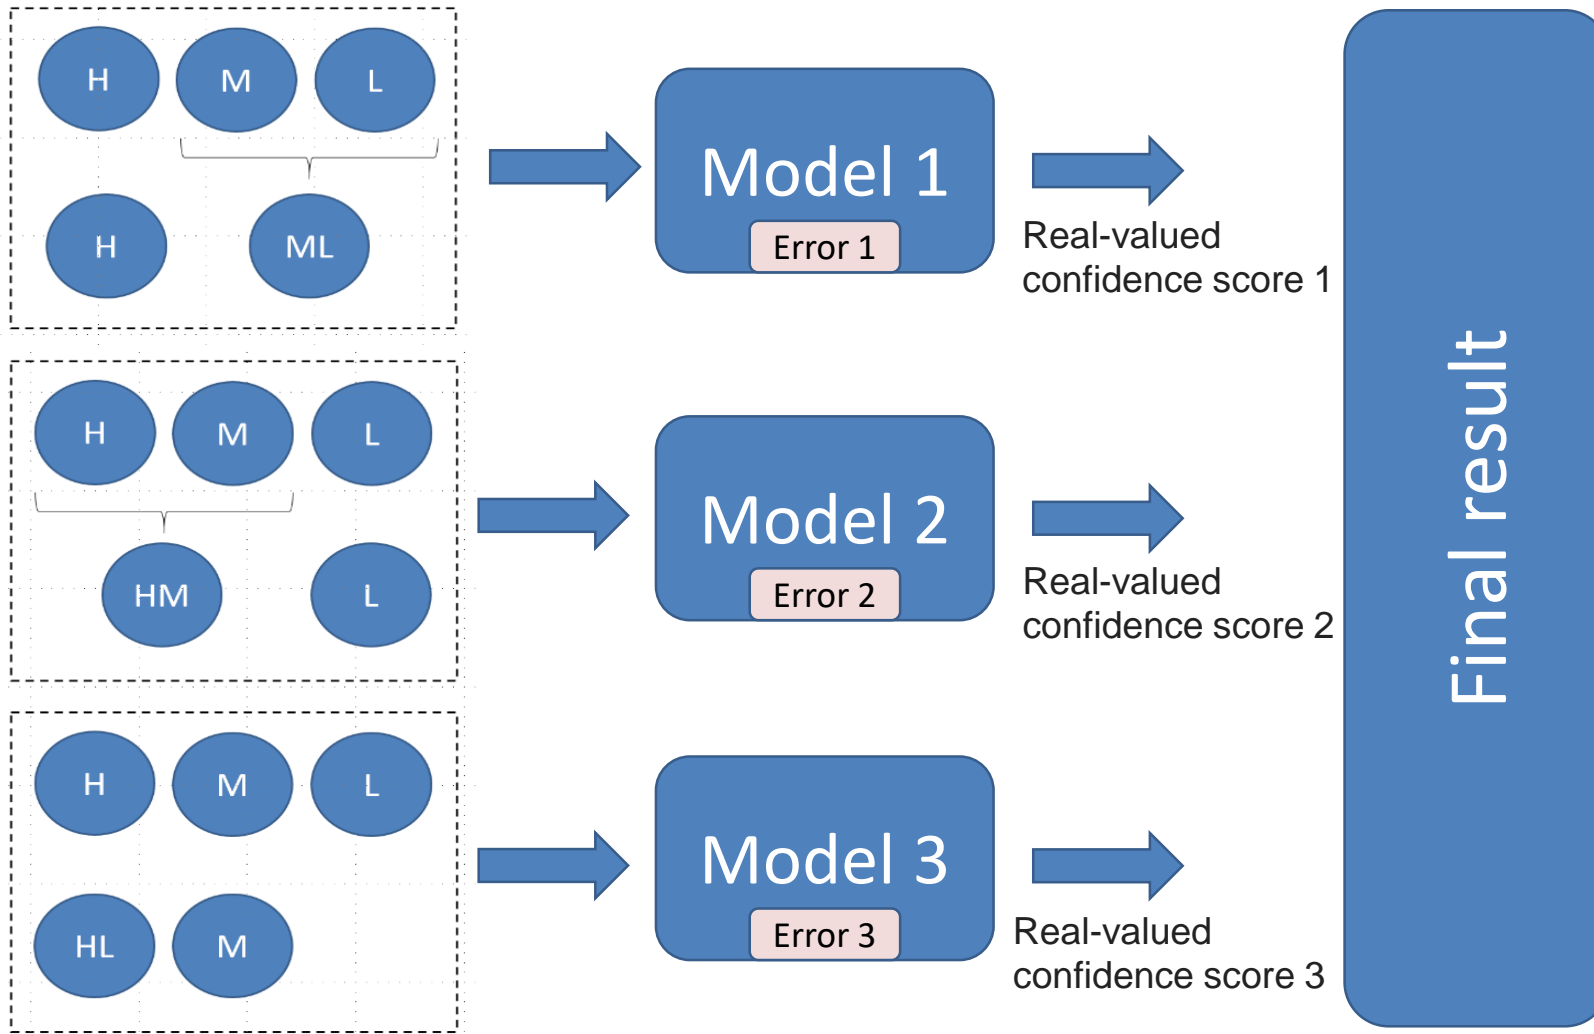


One vs. one approach



Multiclass to binary classification

Training stage



Training a single classifier per class



QUEEN'S
UNIVERSITY
BELFAST

SCHOOL OF
ELECTRONICS,
ELECTRICAL
ENGINEERING AND
COMPUTER SCIENCE

Any Questions?